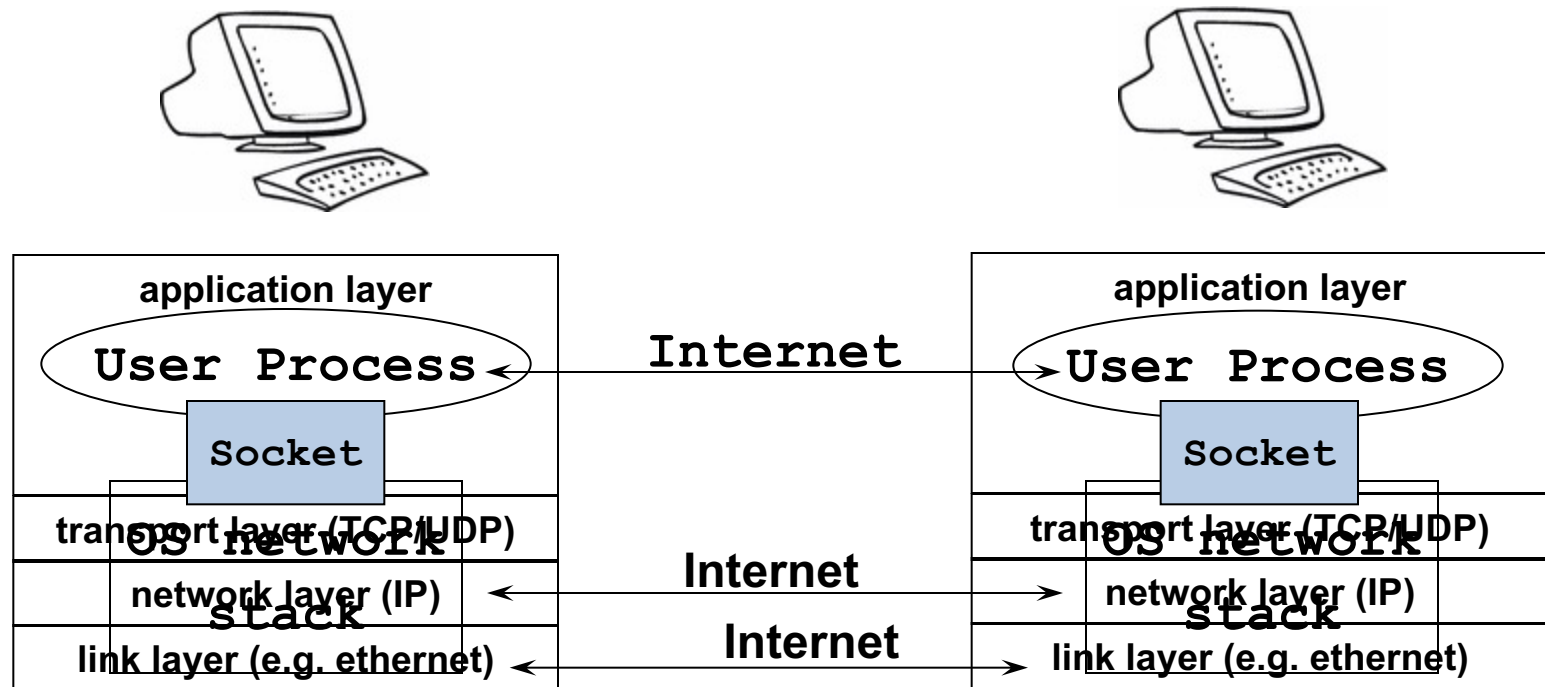


# UNIX Sockets

# Socket and Process Communication



**The interface that the OS provides to its networking subsystem**

# Delivering the Data: Division of Labor

- **Network**

- Deliver data packet to the destination host
- Based on the destination IP address

- **Operating system**

- Deliver data to the destination socket
- Based on the destination port number (e.g., 80)


- **Application**

- Read data from and write data to the socket
- Interpret the data (e.g., render a Web page)



# Socket: End Point of Communication

- Sending message from one process to another
  - Message must traverse the underlying network
- Process sends and receives through a "socket"
  - In essence, the doorway leading in/out of the house
- Socket as an Application Programming Interface
  - Supports the creation of network applications



# Two Types of Application Processes Communication

- **Datagram Socket (UDP)**
  - Collection of messages
  - Best effort
  - Connectionless
- **Stream Socket (TCP)**
  - Stream of bytes
  - Reliable
  - Connection-oriented

# User Datagram Protocol (UDP): Datagram Socket

## UDP

- Single socket to receive messages
  - No guarantee of delivery
- Not necessarily in-order delivery
- Datagram – independent packets
  - Must address each packet

## Postal Mail

- Single mailbox to receive letters
  - Unreliable
- Not necessarily in-order delivery
- Letters sent independently
  - Must address each mail

**Example UDP applications**  
**Multimedia, voice over IP (Skype)**

# Transmission Control Protocol (TCP): Stream Socket

## **TCP**

- **Reliable – guarantee delivery**
- **Byte stream – in-order delivery**
- **Connection-oriented – single socket per connection**
- **Setup connection followed by data transfer**

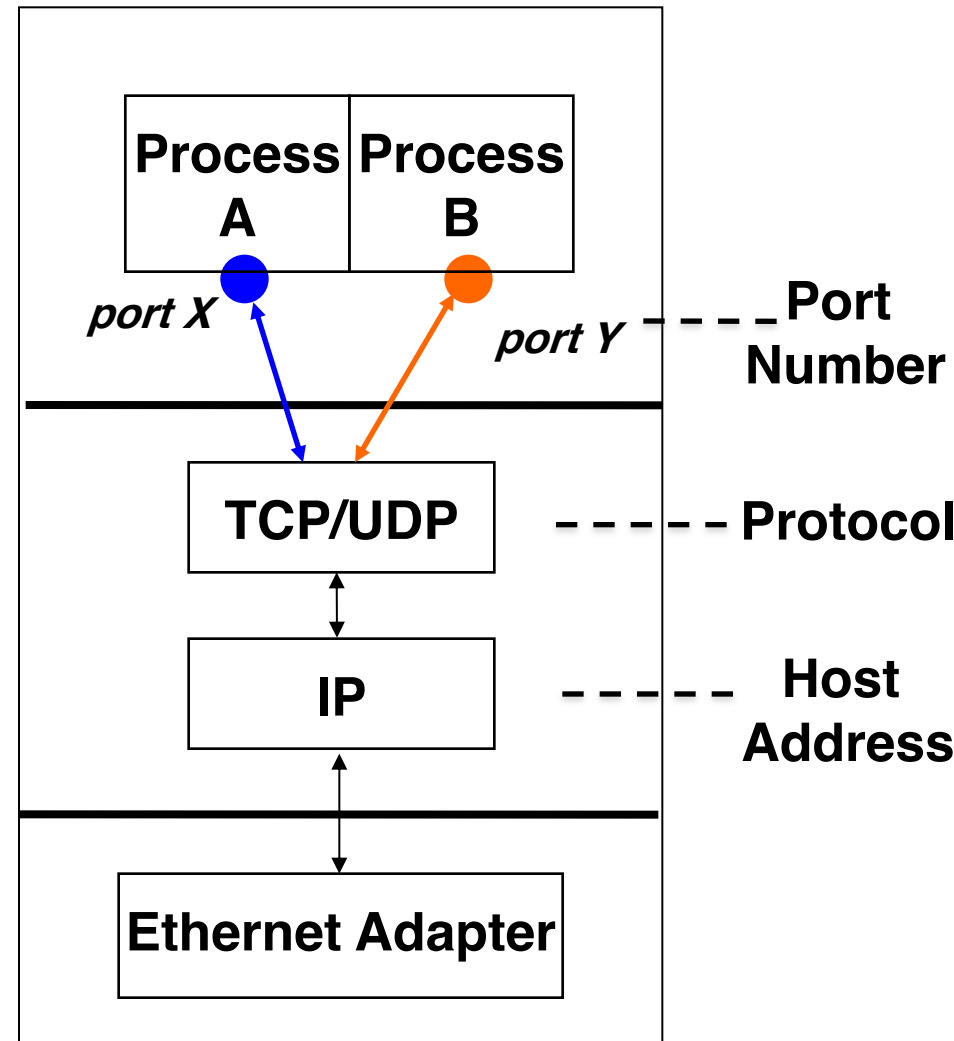
## **Telephone Call**

- **Guaranteed delivery**
- **In-order delivery**
- **Connection-oriented**
- **Setup connection followed by conversation**

**Example TCP applications**  
**Web, Email, Telnet**

# Socket Identification


- **Receiving host**
  - Destination **address** that uniquely identifies host
  - **IP address**: 32-bit quantity
- **Receiving socket**
  - Host may be running many different processes
  - Destination **port** that uniquely identifies socket
  - **Port number**: 16-bits





# Client-Server Communication

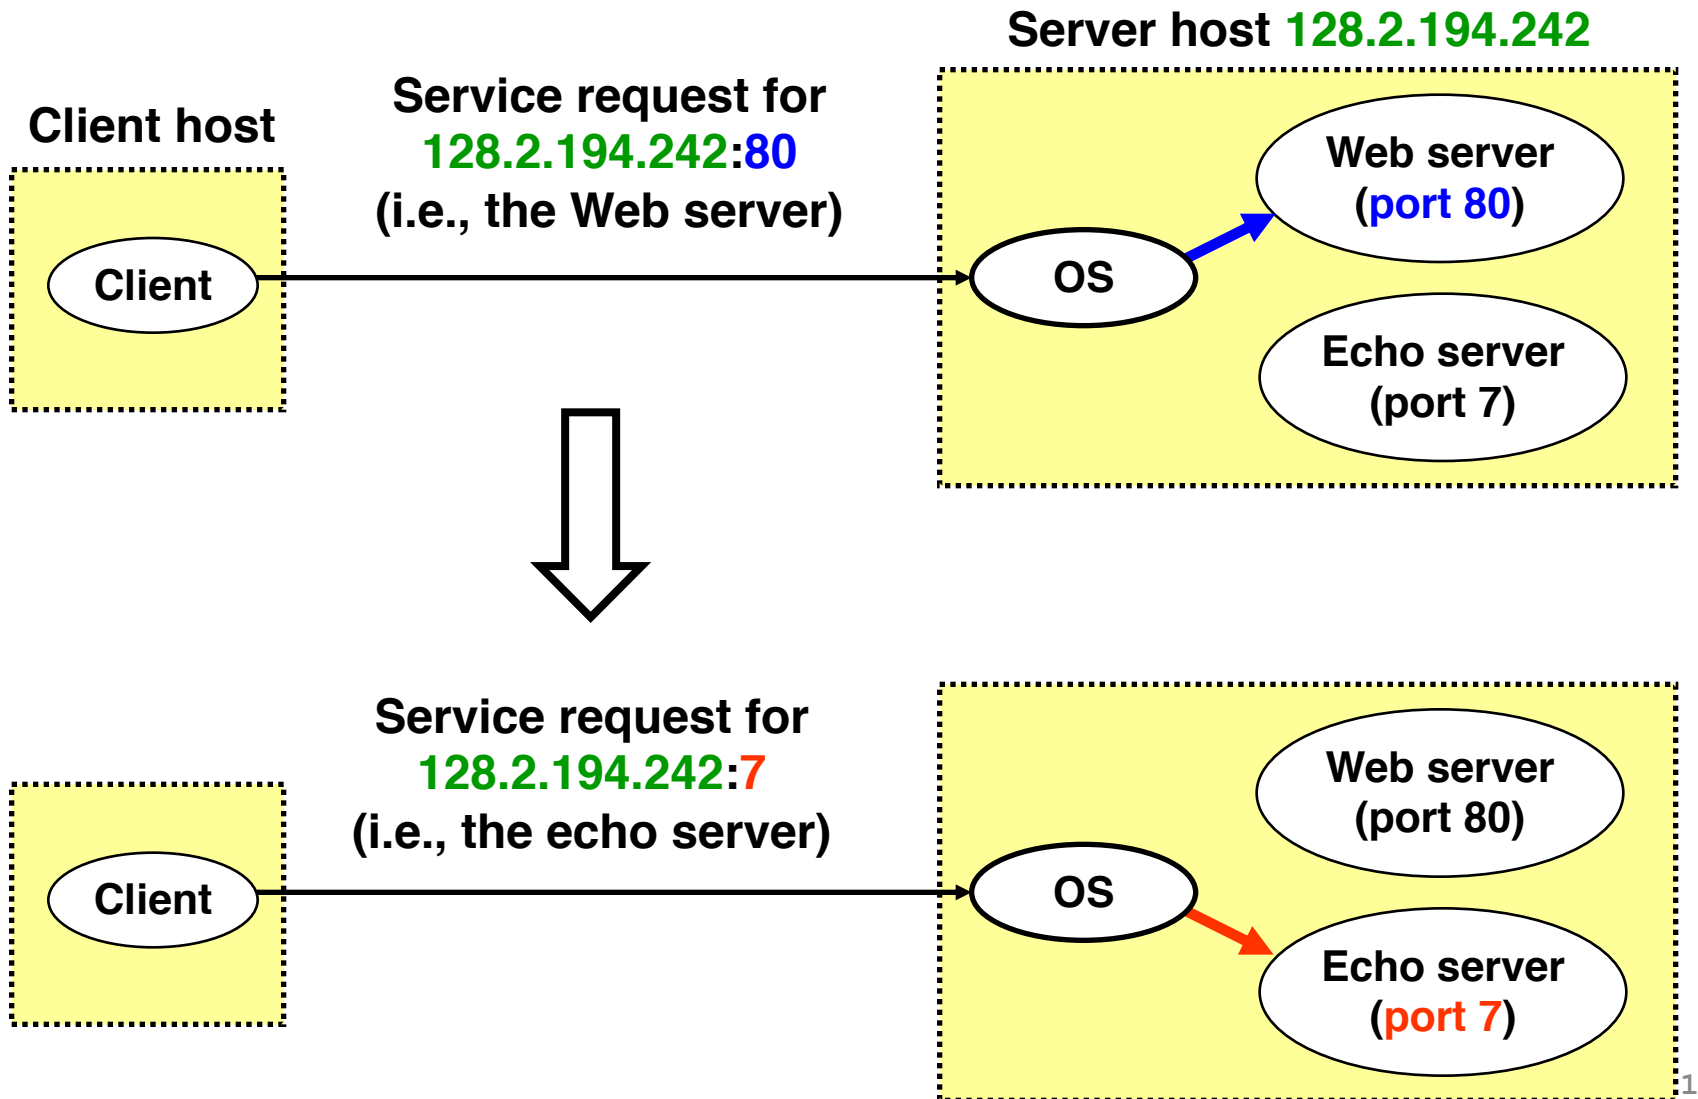
- Client "sometimes on"
  - Initiates a request to the server when interested
  - E.g., Web browser on your laptop or cell phone
  - Doesn't communicate directly with other clients
  - Needs to know server's address
- Server is "always on"
  - Handles services requests from many client hosts
  - E.g., Web server for the [www.cnn.com](http://www.cnn.com) Web site
  - Doesn't initiate contact with the clients
  - Needs fixed, known address



# Knowing What Port Number To Use

- Popular applications have well-known ports
  - E.g., port 80 for Web and port 25 for e-mail
  - See <http://www.iana.org/assignments/port-numbers>
- Well-known vs. ephemeral ports
  - Server has a well-known port (e.g., port 80)
    - Between 0 and 1023 (requires root to use)
  - Client picks an unused ephemeral (i.e., temporary) port
    - Between 1024 and 65535
- "5 tuple" uniquely identifies traffic between hosts
  - Two IP addresses and two port numbers
  - + underlying transport protocol (e.g., TCP or UDP)

# Using Ports to Identify Services



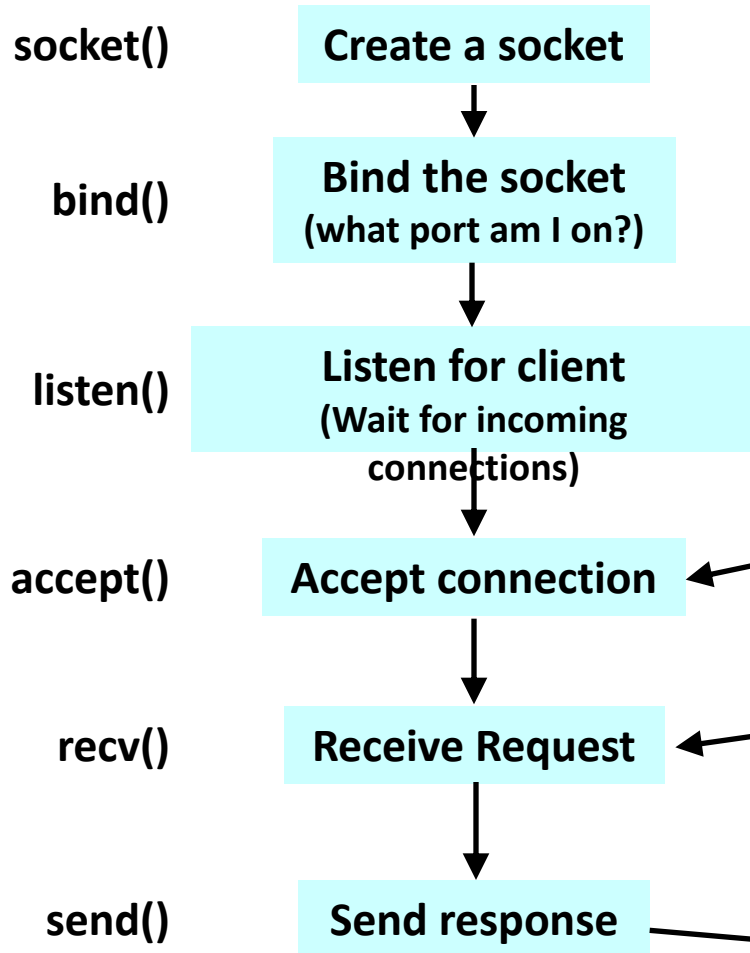
# UNIX Socket API

- In UNIX, everything is like a file
  - All input is like reading a file
  - All output is like writing a file
  - File is represented by an integer file descriptor
- API implemented as system calls
  - E.g., connect, send, recv, close, ...

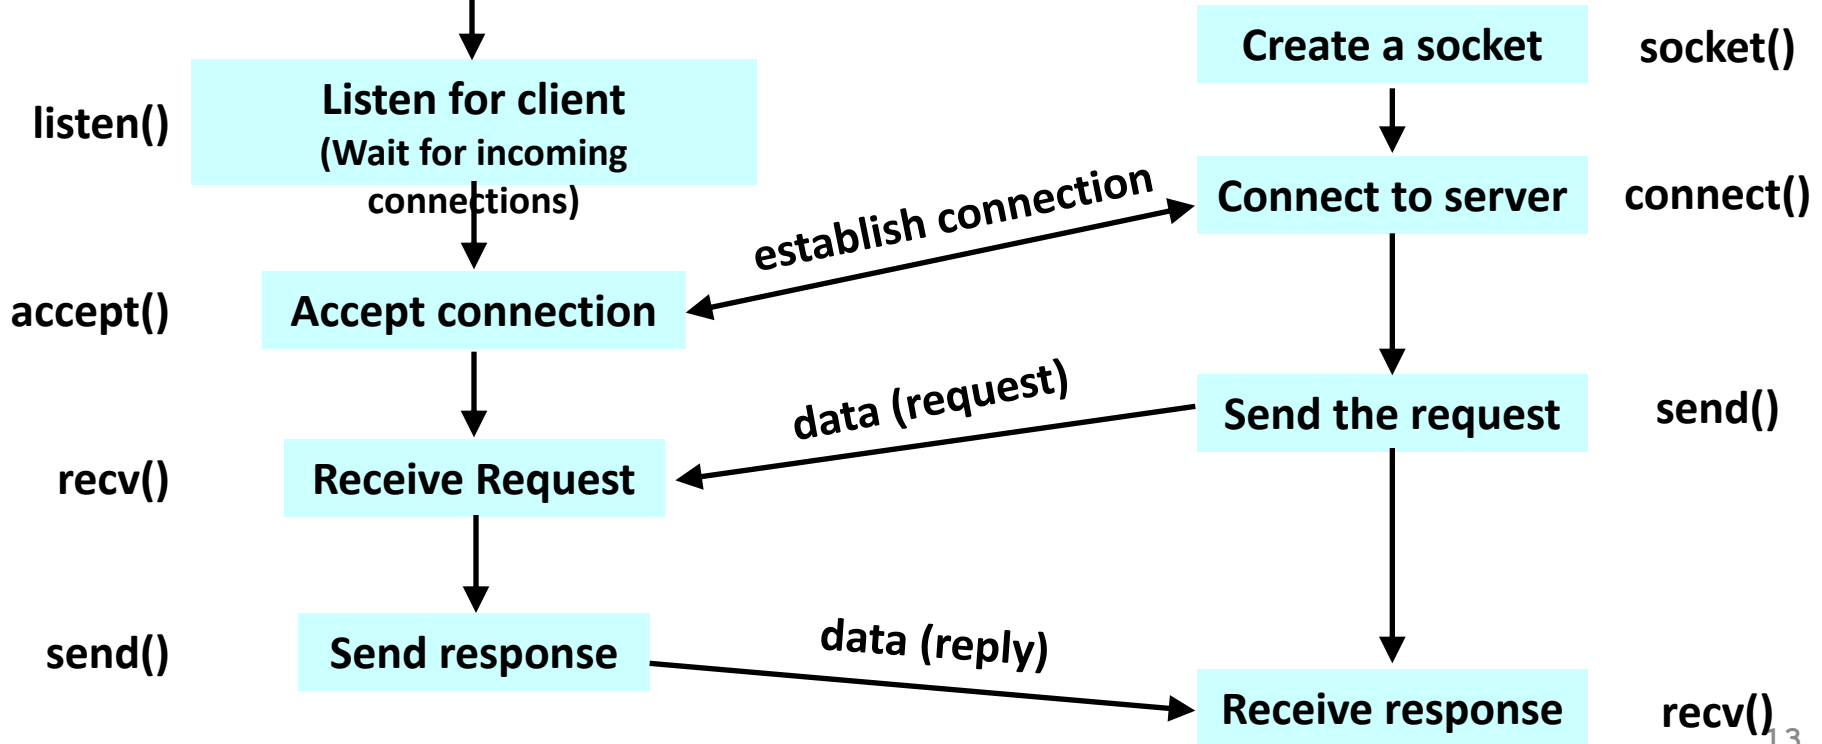
# Client-Server Communication

## Stream Sockets (TCP): Connection-oriented

### Server

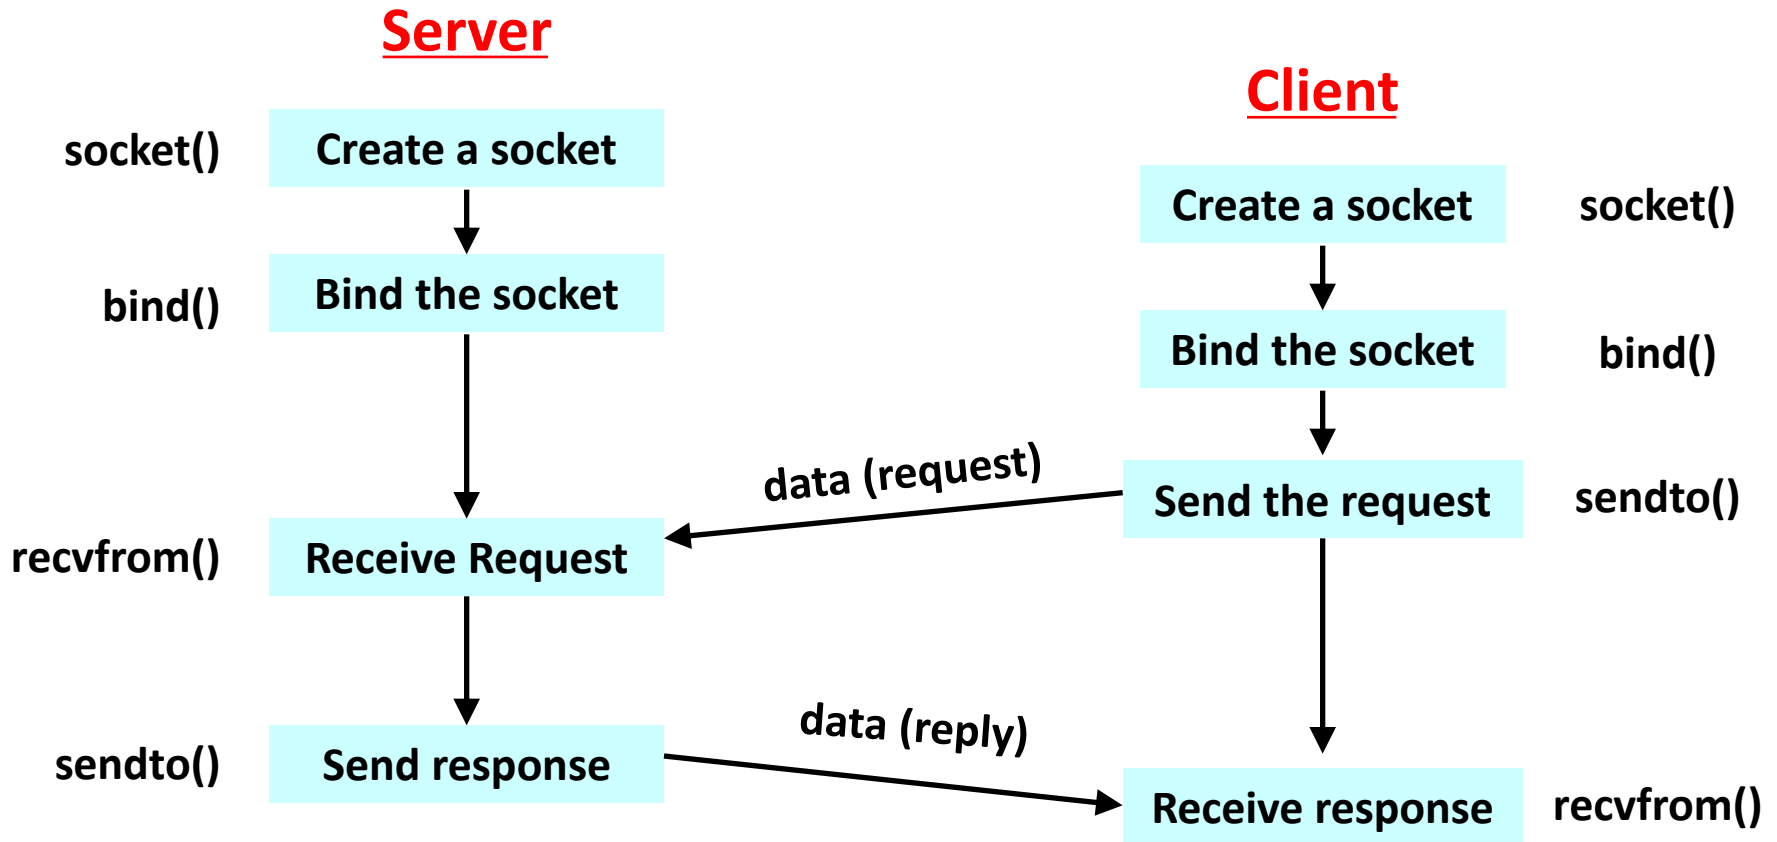


### Client



# Client-Server Communication

## Datagram Sockets (UDP): Connectionless



# Client: Learning Server Address/Port

- Server typically known by name and service
  - E.g., "www.cnn.com" and "http"
- Need to translate into IP address and port #
  - E.g., "64.236.16.20" and "80"
- Get address info with given host name and service
  - `int getaddrinfo( char *node,  
char *service,  
struct addrinfo *hints,  
struct addrinfo **result)`
    - `*node`: host name (e.g., "www.cnn.com") or IP address
    - `*service`: port number or service listed in */etc/services* (e.g. ftp)
    - `hints`: points to a *struct addrinfo* with known information

# Client: Learning Server Address/Port (cont.)

- Data structure to host address information

```
struct addrinfo {  
    int             ai_flags;  
    int             ai_family; //e.g. AF_INET for IPv4  
    int             ai_socktype; //e.g. SOCK_STREAM for TCP  
    int             ai_protocol; //e.g. IPPROTO_TCP  
    size_t          ai_addrlen;  
    char            *ai_canonname;  
    struct sockaddr *ai_addr; // point to sockaddr struct  
    struct addrinfo *ai_next;  
}
```

- Example

```
hints.ai_family = AF_UNSPEC;    // don't care IPv4 or IPv6  
hints.ai_socktype = SOCK_STREAM; // TCP stream sockets  
int status = getaddrinfo("www.cnn.com", "80", &hints, &result);  
// result now points to a linked list of 1 or more addrinfos  
// etc.
```



# Client: Creating a Socket

- **Creating a socket**
  - `int socket(int domain, int type, int protocol)`
  - Returns a file descriptor (or handle) for the socket
- **Domain: protocol family**
  - PF\_INET for IPv4
  - PF\_INET6 for IPv6
- **Type: semantics of the communication**
  - SOCK\_STREAM: reliable byte stream (TCP)
  - SOCK\_DGRAM: message-oriented service (UDP)
- **Protocol: specific protocol**
  - UNSPEC: unspecified
  - (PF\_INET and SOCK\_STREAM already implies TCP)
- **Example**

```
sockfd = socket(    result->ai_family,
                   result->ai_socktype,
                   result->ai_protocol);
```

# Client: Connecting Socket to the Server

- Client contacts the server to establish connection
  - Associate the socket with the server address/port
  - Acquire a local port number (assigned by the OS)
  - Request connection to server, who hopefully accepts
  - connect is **blocking**
- Establishing the connection
  - `int connect( int sockfd,  
                  struct sockaddr *server_address,  
                  socketlen_t addrlen )`
  - Args: socket descriptor, server address, and address size
  - Returns 0 on success, and -1 if an error occurs
  - E.g. `connect( sockfd,  
                  result->ai_addr,  
                  result->ai_addrlen);`

# Client: Sending Data

- Sending data

- `int send( int sockfd, void *msg,  
          size_t len, int flags)`

- Arguments: socket descriptor, pointer to buffer of data to send, and length of the buffer
  - Returns the number of bytes written, and -1 on error
  - send is **blocking**: return only after data is sent
  - Write short messages into a buffer and send once

# Client: Receiving Data

- Receiving data

- `int recv( int sockfd, void *buf, size_t len, int flags)`
- Arguments: socket descriptor, pointer to buffer to place the data, size of the buffer
- Returns the number of characters read (where 0 implies "end of file"), and -1 on error
- Why do you need len? What happens if buf's size < len?
- recv is **blocking**: return only after data is received

# Byte Order

- Network byte order
  - Big Endian
- Host byte order
  - Big Endian (IBM mainframes, Sun SPARC) *or* Little Endian (x86)
- Functions to deal with this
  - htons() & htonl() (host to network short and long)
  - ntohs() & ntohl() (network to host short and long)
- When to worry?
  - putting data onto the wire
  - pulling data off the wire

# Server: Server Preparing its Socket

- Server creates a socket and binds address/port
  - Server creates a socket, just like the client does
  - Server associates the socket with the port number
- Create a socket
  - `int socket( int domain,  
int type, int protocol )`
- Bind socket to the local address and port number
  - `int bind( int sockfd,  
struct sockaddr *my_addr,  
socklen_t addrlen )`

# Server: Allowing Clients to Wait

- Many client requests may arrive
  - Server cannot handle them all at the same time
  - Server could reject the requests, or let them wait
- Define how many connections can be pending
  - `int listen(int sockfd, int backlog)`
  - Arguments: socket descriptor and acceptable backlog
  - Returns a 0 on success, and -1 on error
  - Listen is **non-blocking**: returns immediately
- What if too many clients arrive?
  - Some requests don't get through
  - The Internet makes no promises...
  - And the client can always try again



# Server: Accepting Client Connection

- Now all the server can do is wait...
  - Waits for connection request to arrive
  - **Blocking** until the request arrives
  - And then accepting the new request
- Accept a new connection from a client
  - `int accept( int sockfd,`  
`struct sockaddr *addr,`  
`socketlen_t *addrlen)`
  - Arguments: sockfd, structure that will provide client address and port, and length of the structure
  - Returns descriptor of socket for this new connection





# Client and Server: Cleaning House

- Once the connection is open
  - Both sides can read and write
  - Two unidirectional streams of data
  - In practice, client writes first, and server reads
  - ... then server writes, and client reads, and so on
- Closing down the connection
  - Either side can close the connection
  - ... using the `int close(int sockfd)`
- What about the data still "in flight"
  - Data in flight still reaches the other end
  - So, server can `close()` before client finishes reading

# Server: One Request at a Time?

- Serializing requests is inefficient
  - Server can process just one request at a time
  - All other clients must wait until previous one is done
  - What makes this inefficient?
- May need to time share the server machine
  - Alternate between servicing different requests
    - Do a little work on one request, then switch when you are waiting for some other resource (e.g., reading file from disk)
    - "Nonblocking I/O"
  - Or, use a different process/thread for each request
    - Allow OS to share the CPU(s) across processes
  - Or, some hybrid of these two approaches

# Handle Multiple Clients using fork()

- Steps to handle multiple clients
  - Go to a loop and accept connections using `accept()`
  - After a connection is established, call `fork()` to create a new child process to handle it
  - Go back to listen for another socket in the parent process
  - `close()` when you are done.

```
while (1) {
    fd = accept (srv_fd, (struct sockaddr *) &caddr, &crlen);
    ...
    pid = fork(); children++;
    /* child process to handle request */
    if (pid == 0) {
        /* exit(0) on success, exit(1) on error */
    }
    /* parent process */
    else if (pid > 0) {
        while ((waitpid(-1, &status, WNOHANG)) > 0)
            children--;
        if (children > MAX_PROCESSES)
            ...
    }
    else {
        perror("ERROR on fork");
        exit(1);
    }
}
```

# Helpful Links

- Want to know more?
  - *Beej's guide to network programming*,  
<https://beej.us/guide/bgnet/>
  - Another tutorial at [https://tutorialspoint.com/unix\\_sockets/](https://tutorialspoint.com/unix_sockets/)
  - [http://www.linuxhowtos.org/C\\_C++/socket.htm](http://www.linuxhowtos.org/C_C++/socket.htm)
  - <https://www.geeksforgeeks.org/udp-server-client-implementation-c/>
  - Also please check out the use of Netcat (NC) utility,  
[https://www.thegeekstuff.com/2012/04/nc-command-examples/?utm\\_source=feedburner](https://www.thegeekstuff.com/2012/04/nc-command-examples/?utm_source=feedburner)